

Recommender Systems II

(Algorithm from Netflix Prize)

Approach

- The approach described in the following was an important part of the solution of the winning team in the Netflix \$1,000,000 competition.
- Concluded 21 September 2009
- Prize given to BellKor's Pragmatic Chaos team

Evaluating Rec. Systems

- Consider root-mean-square error (RMSE).

$$\sqrt{\frac{\sum_{u,i} e_{u,i}^2}{N}} = \sqrt{\frac{\sum_{u,i} (r_{u,i} - \hat{r}_{u,i})^2}{N}}$$

- How to minimize it?

We will use gradient descent to minimize each term:

$$e_{u,i}^2 = (r_{u,i} - \hat{r}_{u,i})^2$$

Mean and biases: Baseline

First what will the prediction be?

$$\hat{r}_{u,i} = \mu + b_u + b_i$$

b_u and b_i capture the deviations of user u and item i from the average μ .

E.g., suppose that we want a baseline predictor for the rating of movie **Titanic** by user **Joe**.

Now, say that the average rating over all movies, μ , is 3.7 stars.

Furthermore, **Titanic** is better than an average movie, so it tends to be rated 0.5 stars above the average, i.e. $b_{\text{Titanic}}=0.5$

On the other hand, **Joe** is a critical user, who tends to rate 0.3 stars lower than the average, i.e. $b_{\text{Joe}}=-0.3$.

Thus, the baseline predictor for **Titanic's** rating by **Joe** would be 3.9 stars by calculating $3.7-0.3+0.5$.

Then use gradient descent to minimize $e_{u,i}^2 = (r_{u,i} - \hat{r}_{u,i})^2$ thus finding b_u and b_i .

Mean and biases: Gradient

$$e_{u,i}^2 = (r_{u,i} - \hat{r}_{u,i})^2 = (r_{u,i} - \mu - b_u - b_i)^2$$

We apply regularization to fight overfitting. So, **minimize**:

$$f(b_u, b_i) = (r_{u,i} - \mu - b_u - b_i)^2 + \lambda_1 b_u^2 + \lambda_2 b_i^2$$

Regularization: The last two terms are added to penalize big b_u and b_i values, so that b_u and b_i don't become **too important** for the particular dataset and thus have **overfitting**. It's penalty for being too good for one dataset and suffering for another.

$$\begin{aligned}\frac{\partial f}{\partial b_u} &= \\ -2(r_{u,i} - \mu - b_u - b_i) + 2\lambda_1 b_u &= \\ -2e_{u,i} + 2\lambda_1 b_u &= \\ -2(e_{u,i} - \lambda_1 b_u) &\end{aligned}$$

$$\begin{aligned}\frac{\partial f}{\partial b_i} &= \\ -2(r_{u,i} - \mu - b_u - b_i) + 2\lambda_2 b_i &= \\ -2e_{u,i} + 2\lambda_2 b_i &= \\ -2(e_{u,i} - \lambda_2 b_i) &\end{aligned}$$

Mean and biases: Gradient Descent

Iterate over each $r_{u,i}$, and apply
gradient descent update rules:

$$b_u \leftarrow b_u + \gamma(e_{u,i} - \lambda_1 b_u)$$
$$b_i \leftarrow b_i + \gamma(e_{u,i} - \lambda_2 b_i)$$

The hyperparameters
could be for example:

$$\lambda_1 = 0.02$$

$$\lambda_2 = 0.02$$

$$\gamma = 0.005$$

...grid-search can be used to find better values for
hyperparameters.

30 iterations are often good enough.

Latent Factors

$$\hat{r}_{u,i} = \mu + b_u + b_i + \mathbf{p}_u \cdot \mathbf{q}_i$$

$$\mathbf{p}_{\text{Gus}} = (2, -2)$$

$$\mathbf{q}_{\text{DumbAndDumber}} = (2.2, -1.8)$$

$$\mathbf{p}_{\text{Gus}} \cdot \mathbf{q}_{\text{DumbAndDumber}} = 4.4 + 3.6 = 8$$

Factor is a “characteristic” of users and movies. Both users and movies are mapped to a common “factor” space. The closer a user and a movie are in factor space, the greater their dot-product.

While the figure shows some names for the factors, e.g. “serious/escapist”, “geared toward males/geared toward females”, in general these are latent (hidden) factors for which it is not possible to come up with a name.

We set the number of factors by setting the dimension of \mathbf{p} and \mathbf{q} vectors. Typical dimension is 100.

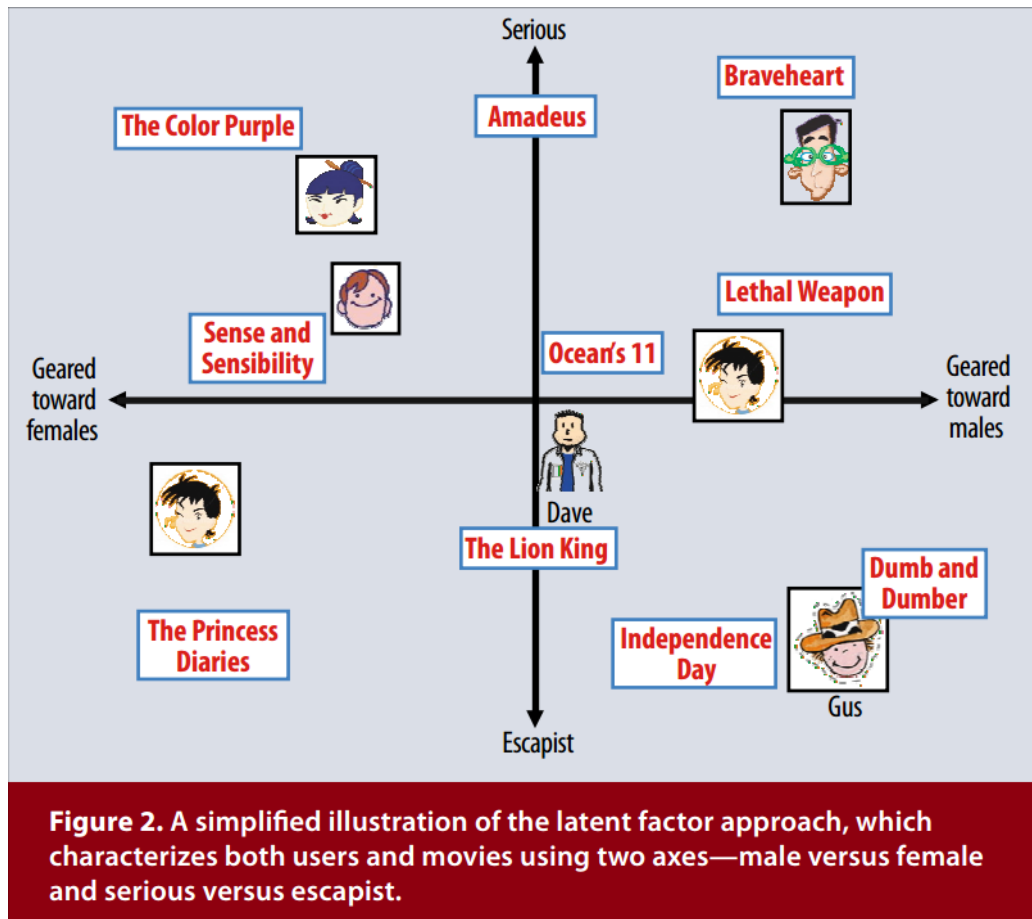


Figure 2. A simplified illustration of the latent factor approach, which characterizes both users and movies using two axes—male versus female and serious versus escapist.

Latent Factors: Gradient

$$e_{u,i}^2 = (r_{u,i} - \hat{r}_{u,i})^2 = (r_{u,i} - \mu - b_u - b_i - \mathbf{p}_u \cdot \mathbf{q}_i)^2$$

We apply regularization to fight overfitting. So, **minimize**:

$$f(b_u, b_i, \mathbf{p}_u, \mathbf{q}_i) = (r_{u,i} - \mu - b_u - b_i - \mathbf{p}_u \cdot \mathbf{q}_i)^2 + \lambda_1 b_u^2 + \lambda_2 b_i^2 + \lambda_3 \|\mathbf{p}_u\|^2 + \lambda_4 \|\mathbf{q}_i\|^2$$

$$\frac{\partial f}{\partial b_u} = -2(e_{u,i} - \lambda_1 b_u)$$

$$\frac{\partial f}{\partial b_i} = -2(e_{u,i} - \lambda_2 b_i)$$

$$\frac{\partial f}{\partial \mathbf{p}_u} = -2(e_{u,i} \mathbf{q}_i - \lambda_3 \mathbf{p}_u)$$

$$\frac{\partial f}{\partial \mathbf{q}_i} = -2(e_{u,i} \mathbf{p}_u - \lambda_4 \mathbf{q}_i)$$

Latent Factors: Gradient Descent

Iterate over each $r_{u,i}$, and apply
gradient descent update rules:

$$b_u \leftarrow b_u + \gamma(e_{u,i} - \lambda_1 b_u)$$

$$b_i \leftarrow b_i + \gamma(e_{u,i} - \lambda_2 b_i)$$

$$\mathbf{p}_u \leftarrow \mathbf{p}_u + \gamma(e_{u,i} \mathbf{q}_i - \lambda_3 \mathbf{p}_u)$$

$$\mathbf{q}_i \leftarrow \mathbf{q}_i + \gamma(e_{u,i} \mathbf{p}_u - \lambda_4 \mathbf{q}_i)$$

The hyperparameters
could be for example:

$$\lambda_1 = \lambda_2 = \lambda_3 = \lambda_4 = 0.02$$

$$\gamma = 0.005$$

...grid-search can be used to find better values for
hyperparameters.

30 iterations are often good enough.

Based on

- Yehuda Koren, Robert Bell. *Advances in Collaborative Filtering*. In Recommender Systems Handbook, Springer 2011, pp 145-186.